

## 로이터 데이터셋 로드하기

In [1]:

```
from keras.datasets import reuters

(train_data, train_labels), (test_data, test_labels) = reuters.load_data(num_words=10000)
```

Using TensorFlow backend.

In [2]:

```
print(len(train_data))
print(len(test_data))
```

8982  
2246

In [3]:

```
train_data[10]
```

Out[3]:

```
[1,
 245,
 273,
 207,
 156,
 53,
 74,
 160,
 26,
 14,
 46,
 296,
 26,
 39,
 74,
 2979,
 3554,
 14,
 46,
 4689,
 4329,
 86,
 61,
 3499,
 4795,
 14,
 61,
 451,
 4329,
 17,
 12]
```

## 로이터 데이터셋을 텍스트로 디코딩하기

In [4]:

```
word_index = reuters.get_word_index()
reverse_word_index = dict([(value, key) for (key, value) in word_index.items()])
decoded_newswire = ' '.join([reverse_word_index.get(i - 3, '?') for i in train_data[0]])
```

In [5]:

```
train_labels[10]
```

Out[5]:

3

## 데이터 인코딩하기

In [6]:

```
import numpy as np

def vectorize_sequences(sequences, dimension=10000):
    results = np.zeros((len(sequences), dimension))
    for i, sequence in enumerate(sequences):
        results[i, sequence] = 1
    return results
```

```
x_train = vectorize_sequences(train_data)
x_test = vectorize_sequences(test_data)
```

In [24]:

```
def to_one_hot(labels, dimension=46):
    results = np.zeros((len(labels), dimension))
    for i, labels in enumerate(labels):
        results[i, labels] = 1
    return results
```

```
one_hot_train_labels = to_one_hot(train_labels)
one_hot_test_labels = to_one_hot(test_labels)
```

In [8]:

```
from keras.utils.np_utils import to_categorical

one_hot_train_labels = to_categorical(train_labels)
one_hot_test_labels = to_categorical(test_labels)
```

In [25]:

```
from keras import models, layers

model = models.Sequential()
model.add(layers.Dense(64, activation='relu', input_shape=(10000,)))
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dense(46, activation='softmax'))
```

In [10]:

```
model.compile(optimizer='rmsprop',  
              loss='categorical_crossentropy',  
              metrics=['accuracy'])
```

In [11]:

```
x_val = x_train[:1000]  
partial_x_train = x_train[1000:]  
y_val = one_hot_train_labels[:1000]  
partial_y_train = one_hot_train_labels[1000:]
```

In [12]:

```
history = model.fit(partial_x_train,  
                    partial_y_train,  
                    epochs=20,  
                    batch_size=512,  
                    validation_data = (x_val, y_val))
```

WARNING:tensorflow:From C:\Users\WJW\Anaconda3\lib\site-packages\tensorflow\python\ops\math\_ops.py:3066: to\_int32 (from tensorflow.python.ops.math\_ops) is deprecated and will be removed in a future version.  
Instructions for updating:  
Use tf.cast instead.  
Train on 7982 samples, validate on 1000 samples

Epoch 1/20  
7982/7982 [=====] - 3s 315us/step - loss: 2.5322 - acc: 0.4955 - val\_loss: 1.7208 - val\_acc: 0.6120

Epoch 2/20  
7982/7982 [=====] - 1s 84us/step - loss: 1.4452 - acc: 0.6879 - val\_loss: 1.3459 - val\_acc: 0.7060

Epoch 3/20  
7982/7982 [=====] - 1s 84us/step - loss: 1.0953 - acc: 0.7651 - val\_loss: 1.1708 - val\_acc: 0.7430

Epoch 4/20  
7982/7982 [=====] - 1s 84us/step - loss: 0.8697 - acc: 0.8165 - val\_loss: 1.0793 - val\_acc: 0.7590

Epoch 5/20  
7982/7982 [=====] - 1s 84us/step - loss: 0.7034 - acc: 0.8472 - val\_loss: 0.9844 - val\_acc: 0.7810

Epoch 6/20  
7982/7982 [=====] - 1s 85us/step - loss: 0.5667 - acc: 0.8802 - val\_loss: 0.9411 - val\_acc: 0.8040

Epoch 7/20  
7982/7982 [=====] - 1s 86us/step - loss: 0.4581 - acc: 0.9048 - val\_loss: 0.9083 - val\_acc: 0.8020

Epoch 8/20  
7982/7982 [=====] - 1s 84us/step - loss: 0.3695 - acc: 0.9231 - val\_loss: 0.9363 - val\_acc: 0.7890

Epoch 9/20  
7982/7982 [=====] - 1s 84us/step - loss: 0.3032 - acc: 0.9315 - val\_loss: 0.8917 - val\_acc: 0.8090

Epoch 10/20  
7982/7982 [=====] - 1s 84us/step - loss: 0.2537 - acc: 0.9414 - val\_loss: 0.9071 - val\_acc: 0.8110

Epoch 11/20  
7982/7982 [=====] - 1s 84us/step - loss: 0.2187 - acc: 0.9471 - val\_loss: 0.9177 - val\_acc: 0.8130

Epoch 12/20  
7982/7982 [=====] - 1s 84us/step - loss: 0.1873 - acc: 0.9508 - val\_loss: 0.9027 - val\_acc: 0.8130

Epoch 13/20  
7982/7982 [=====] - 1s 84us/step - loss: 0.1703 - acc: 0.9521 - val\_loss: 0.9323 - val\_acc: 0.8110

Epoch 14/20  
7982/7982 [=====] - 1s 84us/step - loss: 0.1536 - acc: 0.9554 - val\_loss: 0.9689 - val\_acc: 0.8050

Epoch 15/20  
7982/7982 [=====] - 1s 85us/step - loss: 0.1390 - acc: 0.9560 - val\_loss: 0.9686 - val\_acc: 0.8150

Epoch 16/20  
7982/7982 [=====] - 1s 85us/step - loss: 0.1313 - acc: 0.9560 - val\_loss: 1.0220 - val\_acc: 0.8060

Epoch 17/20  
7982/7982 [=====] - 1s 84us/step - loss: 0.1217 - acc: 0.9579 - val\_loss: 1.0254 - val\_acc: 0.7970

Epoch 18/20  
7982/7982 [=====] - 1s 85us/step - loss: 0.1198 - acc: 0.9582 - val\_loss: 1.0430 - val\_acc: 0.8060

Epoch 19/20

7982/7982 [=====] - 1s 84us/step - loss: 0.1138 - acc: 0.9597 - val\_loss: 1.0955 - val\_acc: 0.7970  
Epoch 20/20  
7982/7982 [=====] - 1s 84us/step - loss: 0.1111 - acc: 0.9593 - val\_loss: 1.0674 - val\_acc: 0.8020

In [14]:


```
import matplotlib.pyplot as plt

loss = history.history['loss']
val_loss = history.history['val_loss']

epochs = range(1, len(loss) + 1)

plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.show()
```




In [15]:

```
plt.clf()
acc = history.history['acc']
val_acc = history.history['val_acc']

plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

plt.show()
```



In [16]:

```
from keras import models, layers

model = models.Sequential()
model.add(layers.Dense(64, activation='relu', input_shape=(10000,)))
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dense(46, activation='sigmoid'))

model.compile(optimizer='rmsprop',
              loss='categorical_crossentropy',
              metrics=['accuracy'])

model.fit(partial_x_train,
          partial_y_train,
          epochs=9,
          batch_size=512,
          validation_data = (x_val, y_val))

results = model.evaluate(x_test, one_hot_test_labels)
```

Train on 7982 samples, validate on 1000 samples

```
Epoch 1/9
7982/7982 [=====] - 1s 115us/step - loss: 3.3043 - acc: 0.4400 - val_loss: 2.8549 - val_acc: 0.5160
Epoch 2/9
7982/7982 [=====] - 1s 85us/step - loss: 2.4160 - acc: 0.4456 - val_loss: 2.0822 - val_acc: 0.3550
Epoch 3/9
7982/7982 [=====] - 1s 85us/step - loss: 1.7412 - acc: 0.3529 - val_loss: 1.6212 - val_acc: 0.3570
Epoch 4/9
7982/7982 [=====] - 1s 84us/step - loss: 1.3375 - acc: 0.4053 - val_loss: 1.3278 - val_acc: 0.6640
Epoch 5/9
7982/7982 [=====] - 1s 84us/step - loss: 0.8727 - acc: 0.8006 - val_loss: 1.0503 - val_acc: 0.7780
Epoch 6/9
7982/7982 [=====] - 1s 85us/step - loss: 0.6391 - acc: 0.8667 - val_loss: 0.9283 - val_acc: 0.8140
Epoch 7/9
7982/7982 [=====] - 1s 85us/step - loss: 0.5074 - acc: 0.8985 - val_loss: 0.9482 - val_acc: 0.8230
Epoch 8/9
7982/7982 [=====] - 1s 86us/step - loss: 0.4239 - acc: 0.9149 - val_loss: 0.9127 - val_acc: 0.8220
Epoch 9/9
7982/7982 [=====] - 1s 86us/step - loss: 0.3549 - acc: 0.9243 - val_loss: 0.9706 - val_acc: 0.8090
2246/2246 [=====] - 0s 119us/step
```

In [17]:

```
results
```

Out[17]:

```
[1.0321700950870532, 0.7845057881207521]
```

In [22]:

```
import copy

test_labels_copy = copy.copy(test_labels)
np.random.shuffle(test_labels_copy)
hits_array = np.array(test_labels) == np.array(test_labels_copy)
float(np.sum(hits_array)) / len(test_labels)
```

Out[22]:

0.19056099732858414

In [18]:

```
predictions = model.predict(x_test)

# predictions[0].shape
# np.sum(predictions[0])
# np.argmax(predictions[0])
```

In [19]:

```
np.sum(predictions[0])
```

Out[19]:

0.0005772114

In [20]:

```
from keras import models, layers

model = models.Sequential()
model.add(layers.Dense(128, activation='relu', input_shape=(10000,)))
model.add(layers.Dense(128, activation='relu'))
model.add(layers.Dense(46, activation='sigmoid'))

model.compile(optimizer='rmsprop',
              loss='categorical_crossentropy',
              metrics=['accuracy'])

model.fit(partial_x_train,
          partial_y_train,
          epochs=9,
          batch_size=512,
          validation_data = (x_val, y_val))
results = model.evaluate(x_test, one_hot_test_labels)
```

Train on 7982 samples, validate on 1000 samples

Epoch 1/9

7982/7982 [=====] - 1s 125us/step - loss: 2.9912 - acc: 0.3419 - val\_loss: 2.2328 - val\_acc: 0.3540

Epoch 2/9

7982/7982 [=====] - 1s 85us/step - loss: 1.8512 - acc: 0.3514 - val\_loss: 1.6510 - val\_acc: 0.3540

Epoch 3/9

7982/7982 [=====] - 1s 85us/step - loss: 1.3002 - acc: 0.3578 - val\_loss: 1.3110 - val\_acc: 0.4420

Epoch 4/9

7982/7982 [=====] - 1s 85us/step - loss: 0.7868 - acc: 0.7844 - val\_loss: 0.9464 - val\_acc: 0.7920

Epoch 5/9

7982/7982 [=====] - 1s 86us/step - loss: 0.5124 - acc: 0.8903 - val\_loss: 0.8999 - val\_acc: 0.8210

Epoch 6/9

7982/7982 [=====] - 1s 86us/step - loss: 0.3969 - acc: 0.9163 - val\_loss: 0.8966 - val\_acc: 0.8210

Epoch 7/9

7982/7982 [=====] - 1s 87us/step - loss: 0.3018 - acc: 0.9347 - val\_loss: 0.9494 - val\_acc: 0.8050

Epoch 8/9

7982/7982 [=====] - 1s 85us/step - loss: 0.2624 - acc: 0.9409 - val\_loss: 0.8960 - val\_acc: 0.8230

Epoch 9/9

7982/7982 [=====] - 1s 85us/step - loss: 1.4963 - acc: 0.6119 - val\_loss: 3.8286 - val\_acc: 0.0060

2246/2246 [=====] - 0s 106us/step

In [23]:

```
results
```

Out[23]:

[3.828641414005419, 0.005342831700801425]